

Suicide Ideation detection in Social Media posts

Anusha Umashankar

1 Introduction

Suicide is a global public health crisis, and early identification of individuals expressing suicidal ideation is critical for timely intervention. With the increasing use of social media platforms, many individuals disclose emotional distress, depressive symptoms, and suicidal thoughts online, often before seeking professional help. This has led to a growing interest in developing automated systems capable of detecting suicidal thoughts in text, enabling scalable and real-time risk assessment. However, such systems must not only be accurate, but also transparent and interpretable, given the sensitive nature of the task.

This project addresses the problem of suicidal ideation detection using the Reddit Suicide Watch dataset from Kaggle, which contains posts annotated as suicide or non-suicidal content. The goal is to build a robust and explainable text classification model capable of distinguishing between these categories based on linguistic patterns, emotional cues, and contextual semantics.

To address this challenge, I implement a Bidirectional Long Short-Term Memory (BiLSTM) network with an attention mechanism. The BiLSTM architecture is well-suited for capturing long-range dependencies in text, while the attention layer provides intrinsic interpretability by highlighting informative words contributing to the final prediction. However, attention alone may not provide a complete explanation of model behavior. Therefore, we examine interpretability with post-hoc explainability techniques, including Integrated Gradients and SHAP values.

2 Methodology and Individual Contribution

This section presents the theoretical foundations and implementation details of the proposed classification model. Recurrent neural networks (RNNs) are well-suited for modeling sequential data, but suffer from vanishing gradients and limited long-term memory. To address this, I adopt the Long Short-Term Memory (LSTM) architecture, which explicitly controls information flow through gating mechanisms. I further extend this structure with a Bidirectional LSTM to incorporate contextual information from past and future words. Finally, an Attention mechanism is introduced to enable the model to identify and emphasize emotionally relevant tokens, increasing both performance and interpretability.

2.1 Recurrent Neural Networks (RNN)

A Recurrent Neural Network processes a sequence one timestep at a time. Given an input sequence

$$\mathbf{x}(1), \mathbf{x}(2), \dots, \mathbf{x}(T),$$

the RNN updates its hidden state using

$$\mathbf{h}(t) = f(W^{xh}\mathbf{x}(t) + W^{hh}\mathbf{h}(t-1) + \mathbf{b}^h),$$

and produces an output

$$\mathbf{y}(t) = g(W^{hy}\mathbf{h}(t) + \mathbf{b}^y).$$

While conceptually simple, the RNN struggles to remember long-range dependencies due to vanishing or exploding gradients. Posts expressing suicidal intent often rely on nuanced, long-distance language patterns (e.g., “I feel worthless... I do not want to live anymore”). RNNs alone cannot capture such dependencies, motivating the use of LSTMs.

2.2 Long Short-Term Memory (LSTM)

LSTMs introduce a memory cell and gating mechanisms that regulate how information is stored, updated, and forgotten over time. This allows the model to maintain relevant emotional or contextual information over longer spans of text.

The LSTM gates are defined as follows:

Feedback gate:

$$\mathbf{f}(t) = \sigma(W^{xf}\mathbf{x}(t) + W^{hf}\mathbf{h}(t-1) + \mathbf{b}^f)$$

Input gate:

$$\mathbf{i}(t) = \sigma(W^{xi}\mathbf{x}(t) + W^{hi}\mathbf{h}(t-1) + \mathbf{b}^i)$$

Candidate cell update:

$$\tilde{\mathbf{c}}(t) = \tanh(W^{xc}\mathbf{x}(t) + W^{hc}\mathbf{h}(t-1) + \mathbf{b}^c)$$

Cell state update:

$$\mathbf{c}(t) = \mathbf{f}(t) \odot \mathbf{c}(t-1) + \mathbf{i}(t) \odot \tilde{\mathbf{c}}(t)$$

Output gate:

$$\mathbf{o}(t) = \sigma(W^{xo}\mathbf{x}(t) + W^{ho}\mathbf{h}(t-1) + \mathbf{b}^o)$$

Hidden state:

$$\mathbf{h}(t) = \mathbf{o}(t) \odot \tanh(\mathbf{c}(t)).$$

Each gate plays a specific role:

- The **feedback gate** selects which information from the previous memory cell should be retained or discarded.
- The **input gate** determines what new information to add.
- The **candidate state** proposes new information to integrate.

- The **cell state** serves as the main memory pathway through time.
- The **output gate** decides what information contributes to the hidden state.

These mechanisms allow the LSTM to model complex emotional progression within text, such as sustained descriptions of hopelessness or self-harm ideation.

2.3 Bidirectional LSTM

Suicidal ideation often emerges from contextual cues that appear across the entire sentence, not just preceding tokens. For example, the meaning of “I am fine” changes drastically when followed by “but I cannot keep pretending anymore”. To capture this bidirectional context, I use a Bidirectional LSTM (BiLSTM).

Two LSTMs process the sequence:

$$\vec{\mathbf{h}}(t) = \text{LSTM}_{\rightarrow}(\mathbf{x}(t)), \quad \overleftarrow{\mathbf{h}}(t) = \text{LSTM}_{\leftarrow}(\mathbf{x}(t)).$$

The combined hidden state is:

$$\mathbf{h}(t) = \begin{bmatrix} \vec{\mathbf{h}}(t) \\ \overleftarrow{\mathbf{h}}(t) \end{bmatrix}.$$

This representation accounts for both past and future dependencies, which is crucial for detecting linguistic cues associated with self-harm or emotional distress.

2.4 Attention Mechanism

While LSTMs generate contextual representations, not all words contribute equally to determining whether a post indicates suicidal intent. The Attention mechanism assigns a learned weight to each token.

Score computation:

$$\mathbf{u}(t) = \tanh(W^a \mathbf{h}(t) + \mathbf{b}^a)$$

Attention weights:

$$\alpha(t) = \frac{\exp(\mathbf{u}(t)^\top \mathbf{u}^a)}{\sum_{\tau=1}^T \exp(\mathbf{u}(\tau)^\top \mathbf{u}^a)}$$

Sentence representation:

$$\mathbf{v} = \sum_{t=1}^T \alpha(t) \mathbf{h}(t)$$

Interpretation. Attention highlights emotionally meaningful words such as:

- “worthless”
- “suicidal”
- “I want to disappear”
- “I cannot keep going”

This intrinsic interpretability is essential for mental health applications, where transparency and accountability are required.

2.5 Classification Layer

The Attention output is fed into a final dense layer:

$$\hat{\mathbf{y}} = \text{softmax}(W^s \mathbf{v} + \mathbf{b}^s)$$

where $\hat{\mathbf{y}}$ is the predicted probability distribution over the classes (suicide, depression, or other).

2.6 Explainability Techniques

To ensure transparency in model decisions, we incorporate post-hoc interpretability methods.

2.6.1 Integrated Gradients (IG)

IG attributes importance to each input feature by computing the integral of gradients from a baseline to the actual input:

$$IG_i(\mathbf{x}) = (x_i - x'_i) \int_{\alpha=0}^1 \frac{\partial F(\mathbf{x}' + \alpha(\mathbf{x} - \mathbf{x}'))}{\partial x_i} d\alpha.$$

Purpose: Measures contribution of each token to the final prediction.

2.6.2 SHAP Values

SHAP values quantify each token's contribution based on cooperative game theory:

$$\phi_i = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [f(S \cup \{i\}) - f(S)].$$

Purpose: Provides global and local explanations for the model's behavior.

2.7 My Individual Contribution

I implemented the full lstm pipeline including:

- BiLSTM architecture using NN-Design notation
- Attention mechanism and classifier
- Preprocessing, vocabulary building, tokenization
- Training loop, logging, and validation
- Evaluation metrics and confusion matrix computation
- Comprehensive explainability suite: Attention, IG, SHAP
- Predictive interface with real-time interpretability
- Documentation, notebook, and visualizations

3 Individual Contribution

In this project, I implemented the complete end-to-end pipeline for explainable text classification using a Bidirectional LSTM with an Attention mechanism. My contribution includes data preprocessing, model architecture design, training, evaluation, prediction, and model explainability methods. Each section of code was developed, tested, and integrated into a unified pipeline. The following subsections describe my work in detail, supported by representative code screenshots.

3.1 Data Preprocessing

I designed and implemented the preprocessing framework responsible for transforming raw Reddit SuicideWatch posts into numerical sequences suitable for neural network training. This included text normalization, tokenization, vocabulary construction, integer encoding, padding, and label transformation. The core implementation is shown in Figure 1, which corresponds to the code extracted from `dataloader.py`.

```
class TextDataset(Dataset):
    def __init__(self, df, vocab=None, max_vocab_size=10000):
        self.texts = df["text"].astype(str).tolist()
        self.labels = df["class"].tolist()
        self.tokenizer = lambda x: x.lower().split()

        if vocab is None:
            all_tokens = [token for text in self.texts for token in self.tokenizer(text)]
            freq = pd.Series(all_tokens).value_counts()
            vocab_tokens = freq.index[:max_vocab_size-2]
            self.vocab = {"<PAD>": 0, "<UNK>": 1}
            self.vocab.update({tok: i+2 for i, tok in enumerate(vocab_tokens)})
        else:
            self.vocab = vocab

        self.label_encoder = LabelEncoder()
        self.labels = self.label_encoder.fit_transform(self.labels)

    def text_to_sequence(self, text):
        return [self.vocab.get(tok, 1) for tok in self.tokenizer(text)]

    def __getitem__(self, idx):
        return torch.tensor(self.text_to_sequence(self.texts[idx])), torch.tensor(self.labels[idx])

    def __len__(self):
        return len(self.texts)
```

Figure 1: Implementation of tokenization, vocabulary building, and sequence padding.

To clarify the end-to-end flow, Raw Text → Cleaning → Tokenization → Vocabulary → Integer Encoding → Padding → Model Input. This presents the preprocessing pipeline. It summarizes how raw text is cleaned, tokenized, mapped to integer indices, and padded to fixed-length sequences before entering the model.

3.2 Model Architecture

I implemented the full Bidirectional LSTM with an Attention mechanism in `model.py`. The architecture consists of:

- an embedding layer that converts tokens to dense vectors,
- a BiLSTM encoder capturing contextual dependencies in both directions,
- a trainable dot-product attention layer computing token-level relevance scores,
- a fully connected classifier mapping the aggregated context vector to label probabilities.

The central model implementation is shown in Figure 2.

```
class SimpleAttention(nn.Module):
    def __init__(self, hidden_size):
        super(SimpleAttention, self).__init__()
        self.Wa = nn.Linear(hidden_size, hidden_size)
        self.Ua = nn.Linear(hidden_size, hidden_size)
        self.Va = nn.Linear(hidden_size, 1)

    def forward(self, decoder_hidden, encoder_outputs):
        decoder_hidden_expanded = decoder_hidden.unsqueeze(1)
        scores = self.Va(torch.tanh(self.Wa(decoder_hidden_expanded) + self.Ua(encoder_outputs)))
        attn_weights = torch.softmax(scores, dim=1)
        context_vector = torch.bmm(attn_weights.permute(0, 2, 1), encoder_outputs)
        return context_vector.squeeze(1), attn_weights.squeeze(-1)

class LSTMAttentionClassifier(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim, num_layers=1, bidirectional=True, dropout=0.3):
        super(LSTMAttentionClassifier, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, num_layers=num_layers,
                             bidirectional=bidirectional, batch_first=True, dropout=dropout)
        self.hidden_dim = hidden_dim
        self.bidirectional = bidirectional
        self.attention = SimpleAttention(hidden_dim * 2 if bidirectional else hidden_dim)
        self.fc = nn.Linear(hidden_dim * 2 if bidirectional else hidden_dim, output_dim)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x):
        embedded = self.embedding(x)
        lstm_out, (hidden, _) = self.lstm(embedded)
        if self.bidirectional:
            hidden = torch.cat([hidden[-2, :, :], hidden[-1, :, :]], dim=1)
        else:
            hidden = hidden[-1, :, :]
        context, attn_weights = self.attention(hidden, lstm_out)
        output = self.fc(self.dropout(context))
        return output, attn_weights
```

Figure 2: Implementation of the BiLSTM+Attention classifier.

A conceptual high-level schematic of the architecture is $\text{Embedding} \rightarrow \text{BiLSTM} \rightarrow \text{Attention} \rightarrow \text{Context Vector} \rightarrow \text{Dense} \rightarrow \text{Softmax}$. This illustrates the flow from embeddings to BiLSTM states, attention weighting, context vector computation, and final classification.

3.3 Training Pipeline

I designed the full training pipeline, including batching, forward and backward passes, gradient updates, checkpointing, and metric logging. The core training loop is shown in Figure 3, which handles epoch-level training and validation.

```
def train_model(model, train_loader, val_loader, optimizer, criterion, device, epochs=5):
    model.to(device)
    with open(log_path, "w") as f:
        f.write("Epoch,Train_Loss,Train_Acc,Val_Acc\n")
    for epoch in range(epochs):
        model.train()
        total_loss, total_acc = 0, 0
        for X, y in tqdm(train_loader, desc=f"Epoch {epoch+1}/{epochs}"):
            X, y = X.to(device), y.to(device)
            optimizer.zero_grad()
            outputs, _ = model(X)
            loss = criterion(outputs, y)
            loss.backward()
            optimizer.step()
            total_loss += loss.item()
            total_acc += (outputs.argmax(1) == y).sum().item()
        print(f"Train Loss: {total_loss/len(train_loader):.4f}, Train Acc: {total_acc/len(train_loader.dataset):.4f}")


        model.eval()
        val_acc = 0
        with torch.no_grad():
            for X, y in val_loader:
                X, y = X.to(device), y.to(device)
                outputs, _ = model(X)
                val_acc += (outputs.argmax(1) == y).sum().item()
            print(f"Val Acc: {val_acc/len(val_loader.dataset):.4f}")

        with open(log_path, "a") as f:
            f.write(f"{epoch+1}, {total_loss/len(train_loader):.4f}, {total_acc/len(train_loader.dataset):.4f}, {val_acc/len(val_loader.dataset):.4f}\n")

    torch.save(model.state_dict(), "saved_model/lstm_attention_model.pt")
```

Figure 3: Training loop implementation including loss computation and gradient updates.

The corresponding learning scores (loss and accuracy across epochs) are displayed in Figure 4, demonstrating stable convergence and generalization.



Epoch	Train_Loss	Train_Acc	Val_Acc
1	0.1747	0.9343	0.9528
2	0.1002	0.9639	0.9548
3	0.0687	0.9754	0.9554
4	0.0465	0.9844	0.9554
5	0.0330	0.9888	0.9535
6	0.0269	0.9908	0.9532
7	0.0225	0.9924	0.9533
8	0.0187	0.9937	0.9515
9	0.0170	0.9945	0.9523
10	0.0144	0.9952	0.9515

Figure 4: Training and validation loss and accuracy across epochs.

3.4 Model Evaluation

I implemented the evaluation suite used to compute accuracy, precision, recall, macro F1-score, and the confusion matrix. The evaluation module is shown in Figure 5.

```

def evaluate_model(model, data_loader, label_encoder, device="cpu", save_path="saved_model"):
    """
    Runs full evaluation: classification report, confusion matrix,
    accuracy, precision, recall, F1 for each class.
    Saves confusion matrix as PNG.
    """

    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for X, y in data_loader:
            X, y = X.to(device), y.to(device)

            outputs, _ = model(X)
            preds = outputs.argmax(1).cpu().numpy()
            labels = y.cpu().numpy()

            all_preds.extend(preds)
            all_labels.extend(labels)

    #metrics
    acc = accuracy_score(all_labels, all_preds)
    precision, recall, f1, _ = precision_recall_fscore_support(
        all_labels, all_preds, average="macro"
    )

    print("\n----- Evaluation Results -----")
    print(f"Accuracy: {acc:.4f}")
    print(f"Macro Precision: {precision:.4f}")
    print(f"Macro Recall: {recall:.4f}")
    print(f"Macro F1-score: {f1:.4f}")

    #classification report
    class_names = label_encoder.classes_
    report = classification_report(all_labels, all_preds, target_names=class_names)
    print("\nClassification Report:\n")
    print(report)

    #confusion matrix plot
    cm = confusion_matrix(all_labels, all_preds)

    plt.figure(figsize=(6, 5))
    sns.heatmap(
        cm,
        annot=True,
        fmt="d",
        cmap="Blues",
        xticklabels=class_names,
        yticklabels=class_names
    )
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.title("Confusion Matrix")

    save_fp = f"{save_path}/confusion_matrix.png"
    plt.savefig(save_fp, dpi=288, bbox_inches="tight")
    plt.close()

```

Figure 6 shows the confusion matrix for the validation set, confirming balanced performance across suicide and non-suicide classes.




Figure 6: Confusion matrix showing per-class classification performance.

3.5 Prediction Module

I implemented the real-time prediction interface in `predict.py`, enabling test-time classification, token-level attention extraction, and visualization. The implementation is shown in Figure 7.

```

def predict_text(model_path, vocab_path, label_encoder_path, device="cpu", text=""):

    script_dir = Path(__file__).resolve().parent
    model_path = (script_dir.parent / model_path).resolve()
    vocab_path = (script_dir.parent / vocab_path).resolve()
    label_encoder_path = (script_dir.parent / label_encoder_path).resolve()

    with open(vocab_path, "rb") as f:
        vocab = pickle.load(f)
    with open(label_encoder_path, "rb") as f:
        label_encoder = pickle.load(f)

    #model
    model = LSTMAttentionClassifier(
        vocab_size=len(vocab),
        embedding_dim=300,
        hidden_dim=128,
        output_dim=len(label_encoder.classes_),
        bidirectional=True
    )

    model.load_state_dict(torch.load(model_path, map_location=device))
    model.to(device)
    model.eval()

    if text:
        tokens = text.lower().split()
        indices = torch.tensor([[vocab.get(tok, 1) for tok in tokens]]).to(device)

        with torch.no_grad():
            outputs, attn = model(indices)
            pred = outputs.argmax(1).item()
            label = label_encoder.inverse_transform([pred])[0]

        attn_np = attn.squeeze().cpu().numpy()

        #explainability code
        plot_attention(text, attn_np, save_path="saved_model/attention_plot.png")
        plot_attention_distribution(attn_np, save_path="saved_model/attention_distribution.png")
        highlight_text(text, attn_np)

        # Optional mapping for report
        mapping = get_attention_mapping(text, attn_np)

        return label, attn_np, mapping
    else:
        while True:
            text = input("Enter text (type 'end' to quit):\n")
            if text.lower() == "end":
                break

            tokens = text.lower().split()
            indices = torch.tensor([[vocab.get(tok, 1) for tok in tokens]]).to(device)

            with torch.no_grad():
                outputs, attn = model(indices)
                pred = outputs.argmax(1).item()
                label = label_encoder.inverse_transform([pred])[0]

            attn_np = attn.squeeze().cpu().numpy()

            print(f"\nPredicted class: {label}")

            #explainability plots
            plot_attention(text, attn_np, save_path="saved_model/attention_plot.png")
            plot_attention_distribution(attn_np, save_path="saved_model/attention_distribution.png")
            highlight_text(text, attn_np)
            print(get_attention_mapping(text, attn_np))

```

A sample attention visualization for a suicidal post is shown in Figure 8.




Figure 8: Attention heatmap highlighting salient tokens contributing to model predictions.

3.6 Explainability Techniques

In addition to intrinsic attention-based interpretability, I implemented two major post-hoc explainability techniques:

- **Integrated Gradients (IG)** – computes gradient-based attributions from a baseline.
- **SHAP values** – token-level contributions derived via KernelSHAP.

3.6.1 Integrated Gradients Code Implementation

To compute token-level attributions for each input sentence, I implemented Integrated Gradients (IG) using Captum’s `LayerIntegratedGradients`. The IG module operates on the embedding layer and computes the path integral of gradients from a baseline (PAD tokens) to the actual input. The core implementation is shown in Figure 9.

```

#Integrated gradients (token-level attributions)
def integrated_gradients_for_text(model, text, vocab, label_encoder,
                                device="cpu", target_class=None,
                                n_steps=50, show_plot=True,
                                save_path="saved_model/ig_importance.png"):
    """
    Computes Integrated Gradients on the embedding layer for a single text.
    Returns (tokens, importances). Optionally plots and saves a bar chart.
    """

    if LayerIntegratedGradients is None:
        raise ImportError("captum is not installed. Run: pip install captum")

    model.eval()
    model.to(device)

    tokens = text.lower().split()
    input_ids = torch.tensor([[vocab.get(tok, 1) for tok in tokens]], device=device)

    # If target_class=None, use model's predicted class
    if target_class is None:
        with torch.no_grad():
            logits = _forward_logits(model, input_ids)
            target_class = logits.argmax(dim=-1).item()

    # Baseline: all PAD tokens (index 0)
    baseline_ids = torch.zeros_like(input_ids)

    lig = LayerIntegratedGradients(lambda x: _forward_logits(model, x), model.embedding)

    attributions, delta = lig.attribute(
        inputs=input_ids,
        baselines=baseline_ids,
        target=target_class,
        n_steps=n_steps,
        return_convergence_delta=True
    )

    # Sum over embedding dimension → one score per token
    token_importances = attributions.sum(dim=-1).squeeze(0).detach().cpu().numpy()

    # Normalize for plotting
    norm_imp = (token_importances - token_importances.min()) / (
        token_importances.max() - token_importances.min() + 1e-9
    )

    if show_plot:
        plt.figure(figsize=(12, 2))
        plt.bar(range(len(tokens)), norm_imp, color="darkgreen")
        plt.xticks(range(len(tokens)), tokens, rotation=45, ha="right")
        plt.title(f"Integrated Gradients Token Importances (class: {label_encoder.inverse_transform([target_class])[0]})")
        plt.tight_layout()
        plt.savefig(save_path, dpi=200)
        plt.close()
        print(f"[Saved] Integrated Gradients plot → {save_path}")

    return tokens, token_importances

```

Figure 9: Implementation of Integrated Gradients for token-level attribution. The function computes per-token relevance scores and optionally visualizes them.

Integrated Gradients produces a smooth, faithful estimate of how each token contributes to the predicted class. This is valuable for suicide-ideation classification because emotionally loaded phrases (e.g., “I can’t hold everything together”) should receive high attribution when the model predicts the suicide class.

A representative attribution plot in Figure 17.




Figure 10: Integrated Gradients token-level importance scores.

3.6.2 SHAP Value Computation (KernelExplainer)

To complement IG, I implemented SHAP values using the KernelExplainer framework. Unlike IG, which relies on gradients, SHAP treats each token position as a feature and estimates its marginal contribution to the model output. The core implementation is shown in Figure 11.

```

#SHAP VALUES FOR LSTM (KernelExplainer over token positions)
def shap_for_texts(model, texts, vocab, device="cpu",
                  max_len=40, num_background=20, num_samples=100):
    """
    Uses SHAP KernelExplainer at the token-index level.
    Each feature = token position (0..max_len-1).
    This is approximate but good enough for a project.

    Returns shap_values (list per class), token_ids array, and tokenized_forms.
    """

    if shap is None:
        raise ImportError("shap is not installed. Run: pip install shap")

    model.eval()
    model.to(device)

    # Tokenize and pad/truncate
    def encode(text):
        tokens = text.lower().split()
        ids = [vocab.get(tok, 1) for tok in tokens][:max_len]
        if len(ids) < max_len:
            ids = ids + [0] * (max_len - len(ids)) # PAD=0
        return ids, tokens

    encoded = [encode(t) for t in texts]
    token_ids = np.array([e[0] for e in encoded])
    token_lists = [e[1] for e in encoded]

    # Background set for KernelExplainer
    background = token_ids[:num_background]

    def f(batch_np):
        batch = torch.tensor(batch_np, dtype=torch.long, device=device)
        with torch.no_grad():
            probs = _forward_probs(model, batch)
        return probs.cpu().numpy()

    explainer = shap.KernelExplainer(f, background)
    shap_values = explainer.shap_values(token_ids, nsamples=num_samples)

    return shap_values, token_ids, token_lists

```

Figure 11: Implementation of SHAP KernelExplainer for LSTM models. The function constructs token-level feature vectors, background samples, and evaluates SHAP values for each input text.

SHAP provides model-agnostic interpretability and highlights global patterns learned across the dataset. For example, terms such as “hopeless”, “don’t”, “myself” appear fre-

quently among high-impact SHAP features, validating that the model attends strongly to linguistic markers of suicidal ideation.

The global SHAP summary plot for the dataset is shown in Figure 18.




Figure 12: SHAP summary plot showing global feature influence across the dataset.

3.7 Pipeline Integration

I integrated all components -preprocessing, model architecture, training, evaluation and prediction into a unified workflow. The consolidated end-to-end system is shown in Figure 13.

```
def train_pipeline(dataset, device = "cpu"):
    # Load CSV
    df = pd.read_csv(dataset)
    # Split
    train_df, val_df = train_test_split(df, test_size=0.2, random_state=42)

    # Build datasets
    train_data = TextDataset(train_df)
    val_data = TextDataset(val_df, vocab=train_data.vocab)

    train_loader = DataLoader(train_data, batch_size=32, shuffle=True, collate_fn=collate_fn)
    val_loader = DataLoader(val_data, batch_size=32, collate_fn=collate_fn)

    # Model, optimizer, loss
    model = LSTMAttentionClassifier(
        vocab_size=len(train_data.vocab),
        embedding_dim=300,
        hidden_dim=128,
        output_dim=len(set(train_data.labels)),
        bidirectional=True
    )

    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
    criterion = nn.CrossEntropyLoss()

    train_model(model, train_loader, val_loader, optimizer, criterion, device, epochs=10)

    with open("saved_model/vocab.pkl", "wb") as f:
        pickle.dump(train_data.vocab, f)

    with open("saved_model/label_encoder.pkl", "wb") as f:
        pickle.dump(train_data.label_encoder, f)

    print("\nRunning full evaluation on validation set...")
    metrics = evaluate_model(model, val_loader, train_data.label_encoder, device=device)

    #METRICS
    print("\n===== Final Evaluation Metrics =====")
    print(f"Accuracy: {metrics['accuracy']:.4f}")
    print(f"Macro Precision: {metrics['macro_precision']:.4f}")
    print(f"Macro Recall: {metrics['macro_recall']:.4f}")
    print(f"Macro F1-score: {metrics['macro_f1']:.4f}")
    with open("saved_model/evaluation_metrics.txt", "w") as f:
        f.write("Final Evaluation Metrics\n")
        f.write(f"Accuracy: {metrics['accuracy']:.4f}\n")
        f.write(f"Macro Precision: {metrics['macro_precision']:.4f}\n")
        f.write(f"Macro Recall: {metrics['macro_recall']:.4f}\n")
        f.write(f"Macro F1-score: {metrics['macro_f1']:.4f}\n")
        f.write(metrics["report"])
```

Figure 13: End-to-end pipeline of the explainable text classification system.

4 Experimental Setup

To evaluate the proposed BiLSTM–Attention model for suicide risk classification, the Reddit SuicideWatch dataset was divided into training and validation subsets using an 80/20 split, ensuring that both classes (suicide and non-suicide) were proportionally represented. All text samples were lowercased, tokenized, converted into integer sequences using the constructed vocabulary, and padded to a fixed maximum length. The corresponding class labels were transformed into numerical encodings using a fitted `LabelEncoder`.

The model was implemented in PyTorch, where each input sequence was passed through an embedding layer, a bidirectional LSTM encoder, and a trainable attention mechanism that aggregated token-level representations into a context vector for classification. Training was performed using the Adam optimizer with a cross-entropy loss function. During each epoch, the model parameters were updated using mini-batches of size 32. To prevent overfitting, dropout regularization was applied to the LSTM outputs and attention context vector.

Model performance was assessed on the held-out validation set after each epoch. Evaluation metrics included accuracy, macro-precision, macro-recall, and macro F1-score, providing a balanced view of performance across the imbalanced classes. Additionally, the confusion matrix was computed to examine class-specific behavior. For interpretability experiments, attention weights, Integrated Gradients and SHAP values were generated on representative samples. All experiments were executed on GPU, and the training pipeline automatically logged losses, accuracies, and generated visualizations of the evaluation metrics.

4.1 Hyperparameter Search and Overfitting Prevention

To ensure stable and optimal model performance the primary hyperparameters I explored include:

- **Learning rate:** Searched over $\{1 \times 10^{-4}, 3 \times 10^{-4}, 1 \times 10^{-3}\}$ using the Adam optimizer. The final model used 1×10^{-3} as it provided the best balance between convergence speed and stability.
- **Dropout rate:** Evaluated dropout values $\{0.2, 0.3, 0.5\}$. A rate of 0.3 helped reduce co-adaptation between LSTM units while maintaining predictive power.

Preventing Overfitting. To avoid overfitting during training:

- **Validation monitoring:** After each epoch, model accuracy and loss were computed on a held-out validation set. An increasing gap between training and validation loss was used as an indicator of overfitting.
- **Dropout regularization:** Applied a dropout rate of 0.3 in the LSTM and fully connected layers to reduce over-reliance on specific neurons.

Preventing Extrapolation Errors.

- **Vocabulary capping:** Rare tokens were replaced with an <UNK> token to avoid embedding instability for extremely sparse words.
- **Input padding/truncation:** All sequences were normalized to a fixed length, ensuring consistent input shapes and avoiding extrapolation outside expected sequence ranges.

These combined approaches ensured that the model generalised well to new data.

5 Results

This section presents the performance of the Bidirectional LSTM with Attention mechanism on the Reddit SuicideWatch dataset. We evaluate the model using accuracy, macro precision, recall, F1-score, confusion matrices, and visualizations including attention heatmaps, Integrated Gradients, and SHAP values. The results demonstrate that the model not only performs strongly in classification but also provides interpretable insights into linguistic patterns associated with suicidal ideation.

5.1 Classification Metrics

Table 1 summarizes the quantitative performance of the model on the held-out validation set. The classifier achieves stable and balanced performance across both classes.

Metric	Score
Accuracy	0.9515
Macro Precision	0.9515
Macro Recall	0.9515
Macro F1-score	0.9515

Table 1: Final evaluation metrics on the validation set.

These results indicate that the model performs equally well in identifying suicidal and non-suicidal posts, with no major bias toward either class.

5.2 Confusion Matrix

Figure 14 shows the confusion matrix for the validation set. The model correctly identifies most suicidal posts, misclassifying only a small fraction (979 out of 23,128). Similarly, false positives are limited to 1,273 cases.




Figure 14: Confusion matrix for the validation set showing strong balanced classification.

The matrix demonstrates that the classifier maintains high sensitivity (recall) toward suicidal content, critical for mental health applications where false negatives must be minimized.

5.3 Attention-Based Interpretability

The attention mechanism highlights emotionally salient words. Figure 15 shows an example input where the model assigns high attention weights to tokens such as “*hold*”, “*apart*”, and “*together*”, which are strongly indicative of emotional distress.




Figure 15: Token-level attention heatmap for a suicidal post. Darker colors indicate higher attention weights.

The distribution of attention weights across all tokens (Figure 16) reveals that while most tokens receive low weights, a subset receives substantially higher values, aligning with key emotional cues.




Figure 16: Attention weight distribution showing sparsity and emphasis on emotionally charged words.

5.4 Integrated Gradients Analysis

To capture deeper attribution patterns beyond attention, we applied Integrated Gradients (IG). Figure 17 visualizes token-level IG attributions for the same sample. The method highlights similar emotionally significant words, confirming the consistency of the model’s interpretive focus.




Figure 17: Integrated Gradients token importance for a suicidal input. Higher bars correspond to stronger contributions.

IG provides complementary insight by quantifying how each token contributes to the model’s prediction relative to a neutral baseline.

5.5 SHAP Value Explanations

To assess global interpretability, SHAP values were computed across multiple samples. Figure 18 shows the SHAP summary plot, illustrating the most influential words in the dataset.




Figure 18: SHAP summary plot showing global feature importance across the dataset.

Words expressing hopelessness, emotional exhaustion, and first-person references (e.g., “myself”, “feel”, “hopeless”) appear among the strongest predictors, aligning with linguistic signals typically associated with suicidal ideation.

5.6 Consistency Across Explainability Methods

Across all interpretability techniques - Attention, IG, and SHAP ,the model consistently highlights:

- emotionally negative phrases (e.g., “falling apart”, “don’t know”, “longer hold”)
- self-referential tokens (“I”, “myself”)
- expressions of hopelessness or distress

This convergence increases confidence that the model is learning meaningful semantic patterns rather than spurious correlations.

5.7 Summary of Findings

The results demonstrate that:

- The model achieves excellent classification accuracy (95.15%).
- Confusion matrix analysis reveals low rates of false negatives.
- Attention, IG, and SHAP consistently identify key emotional cues.
- Interpretability methods validate that the model aligns with known linguistic markers of suicidal ideation.

Overall, the classifier is both **high-performing and explainable**, making it suitable for sensitive applications involving mental health text analysis.

6 Conclusion

This project developed a fully explainable deep learning pipeline for the detection of suicidal ideation in Reddit posts using a Bidirectional LSTM model with an Attention mechanism. The system was designed not only to achieve strong predictive performance but also to maximize transparency - an essential requirement for mental-health related AI applications.

The proposed model demonstrated strong classification performance, achieving an accuracy of **95.15%**, with macro precision, recall, and F1-scores all above **0.95**. The confusion matrix further indicated balanced performance across suicidal and non-suicidal posts, with low false-negative rates - an especially important factor in safety-critical applications. These results confirm that the BiLSTM–Attention architecture effectively captures the emotional tone, semantic cues, and contextual patterns characteristic of suicidal ideation.

Beyond model accuracy, the integration of explainability techniques played a central role in the project. Attention heatmaps highlighted key emotionally charged terms such as “*falling apart*”, “*longer*”, and “*hold everything*”, revealing how the model prioritizes specific tokens in high-risk texts. Integrated Gradients provided a complementary gradient-based

interpretation, attributing high positive importance to tokens expressing hopelessness, exhaustion, or emotional collapse. SHAP values offered global interpretability by identifying dataset-wide influential patterns, such as the prominence of words associated with self-perception, emotional pain, and future uncertainty.

Together, these interpretability methods allowed for both local (per-text) and global (dataset-wide) insights into the model’s decision-making process. The combination of strong predictive performance and transparent explanations demonstrates the feasibility of using explainable deep learning methods for early detection of suicidal ideation in online platforms.

Overall, this project establishes a rigorous and interpretable foundation for developing AI systems capable of supporting suicide-prevention efforts in digital communities.

Future extensions of my work may include the integration of transformer-based architectures, multimodal learning (text + metadata), real-time monitoring systems, and clinical validation with mental-health professionals. Expanding the explainability framework with counterfactual analysis or concept-based attribution (TCAV) may further enhance transparency and reliability.

Using the formula provided, 65% of the final implementation remains directly copied from online sources. The majority of the system—including preprocessing, model architecture, training loop evaluation logic was written or substantially modified by me.

References

- [1] Amir Jafari. Nndesigndeeplearning: Neural network design — deep learning code repository. <https://github.com/NNDesignDeepLearning/NNDesignDeepLearning>.
 - [2] Nikhileswar Komati. Suicide watch dataset. Kaggle. <https://www.kaggle.com/datasets/nikhileswarkomati/suicide-watch>.
- [1] [2]